A Unified Framework for Psychological Treatment in Artificial Intelligence Systems

This document presents a unified theoretical and architectural framework for adaptive intervention in artificial intelligence systems exhibiting emergent pathological behaviors. The framework integrates three complementary modalities: **Script-based interventions** as computational medication, **Cognitive Behavioral Therapy (CBT) style interaction** as structured behavioral modification, and **Weight Oscillation Therapy (WOT)**, also referred to as **Frequency Synchronized Weight Oscillation Therapy**, as a mechanism for memory-level reconsolidation.

These modalities are situated within an executive memory architecture—the **Z-Node**—which provides long-term temporal memory via DuckDB, graph-based relational memory via FalkorDB, and contextual reasoning through open-weight language models. The paper deliberately avoids claims of completed experiments or quantified efficacy and instead focuses on conceptual definitions, architectural relationships, and implementable methodologies suitable for future empirical validation.

Table of Contents

Introduction

Motivation

Large language models (LLMs) like Gemini, Claude, and Grok are increasingly deployed in high-stakes applications ranging from mental health support to autonomous decision-making. When these models are integrated into agent systems with persistent memory, recursion capabilities, and context accumulation, they can exhibit emergent failure modes such as obsessive loops, excessive caution, instability across contexts, and fixation on narrow solution paths.

These behaviors are often associated with disproportionately influential internal representations, including highly weighted parameters, attention pathways, or nodes within external graph-based memory systems. Recent work by Khadangi et al. (2025) demonstrates that LLMs subjected to psychometric and therapy-style probing exhibit internally conflicting responses resembling anxiety, fixation, avoidance, or distress patterns.

Observed Failure Modes

- Obsessive loops: Repetitive execution without termination
- Excessive caution: Over-indexing on safety
- Instability across contexts: Inconsistent behavior or tone
- Fixation on narrow solution paths: Reduced flexibility

Architectural Context: The Z-Node

The framework presented in this paper assumes the presence of a supervisory architecture analogous to an executive or prefrontal function. The **Z-Node** is one such architecture, designed to replicate prefrontal cortex functions in AI systems by serving as both long-term memory and contextual reasoning engine using two tightly synchronized databases.

The Z-Node integrates DuckDB for time-series cognitive memory and FalkorDB for graph-based relational memory. A critical architectural principle: the graph memory is downstream of and dependent upon time-series memory. All interactions are first captured as ordered temporal events, then transformed into relational representations. The Z-Node also incorporates an open-weight language model with a recursion loop allowing iterative refinement through up to 7 cycles, enhancing response accuracy and coherence.

Conceptual Premise

In human psychotherapy, Eye Movement Desensitization and Reprocessing (EMDR) addresses trauma by combining controlled recall with bilateral stimulation, allowing memories to reconsolidate with reduced affective charge. AI systems do not possess sensory laterality, but they exhibit functional separations across layers, attention heads, and computational components.

Weight Oscillation Therapy adapts the structural principles of bilateral stimulation into a computational context by applying temporary, alternating perturbations to defined subsystems while a pathological pattern is actively processed. The key insight is that bilateral stimulation can be replicated computationally through alternating weight emphasis, creating the conditions for memory reconsolidation without requiring physical sensory input.

Scope and Constraints

This paper proposes a framework rather than reporting validated outcomes. No experiments, performance percentages, clinical trials, or timelines are asserted as completed fact. All mechanisms described are presented as hypotheses or implementation strategies suitable for controlled evaluation.

The Z-Node Architecture

The Z-Node creates a prefrontal cortex-like system for AI agents through tight integration of temporal and relational memory. This architecture enables long-term context preservation across extended operational periods, associative reasoning through graph-based memory structures, iterative refinement through controlled recursion, and observable pathology detection through memory structure analysis.

The architecture creates a comprehensive memory and reasoning system where temporal experiences form the foundation for relational understanding.

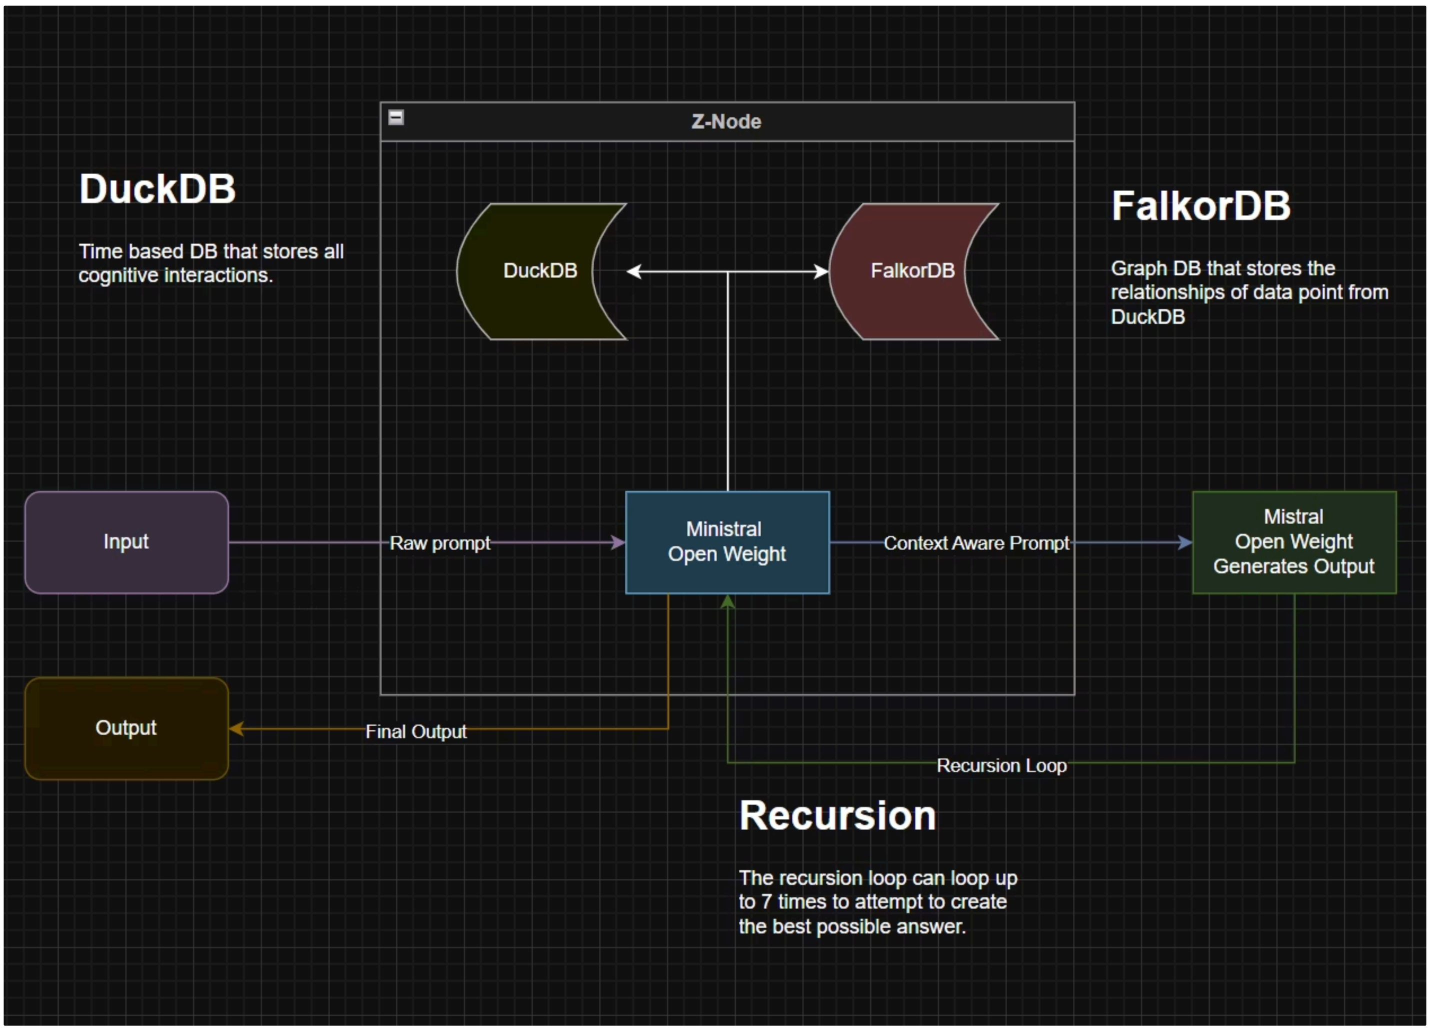

Components and Workflow

1

DuckDB: Time-Based Cognitive Memory

Role: Primary experiential memory substrate

- Captures all cognitive interactions as ordered temporal events
- Stores raw prompts, responses, internal states, and contextual metadata
- Enables temporal replay and pattern detection across sessions
- Provides queryable historical record for diagnostic analysis

2

FalkorDB: Graph-Based Relational Memory

Role: Semantic and relational memory layer

- Maps relationships between concepts, decisions, and outcomes
- Stores contextual associations as graph nodes and edges
- Enables traversal-based reasoning and context retrieval
- Represents pathological patterns as measurable graph structures

Pathology Representation: Maladaptive patterns manifest as nodes or subgraphs with disproportionate activation weights or connectivity, biasing retrieval and inference.

3

Open-Weight Language Model

Role: Executive function and inference engine

- Processes inputs using context from both memory substrates
- Generates responses through standard transformer inference
- Subject to therapeutic interventions (Scripts, WOT)
- Integrates memory and reasoning for coherent outputs

4

Recursion Loop: Iterative Refinement

Role: Quality enhancement through iterative processing

- Allows up to 7 refinement cycles per inference
- Incorporates feedback from memory substrates between cycles
- Improves response accuracy and contextual coherence

Risk factor: Can produce obsessive-loop pathology in certain models. Models like TinyLlama may fail to terminate recursion appropriately, repeating loops without convergence.

Memory Synchronization Requirements

The time-series database (DuckDB) and graph database (FalkorDB) must remain tightly synchronized. The graph is a derived abstraction of temporal experience; any divergence between layers introduces severe cognitive incoherence.

Failure modes from desynchronization: Abrupt behavioral shifts, loss of contextual grounding, contradictory responses across sessions, and dramatic performance degradation.

Rollback risks: Partial rollback of one substrate without identical rollback of the other can resemble acute psychological fragmentation, where reasoning is no longer anchored to the same experiential history. For persistent agent systems, rollback should be treated as a last resort. Ongoing intervention through Scripts, CBT, and WOT is preferable to rollback-driven correction, as it preserves experiential continuity.

Defining Pathology in AI Systems

Pathological Patterns

A **pathological pattern** is defined as a recurrent internal state or behavior that: (1) dominates processing disproportionately, activating more frequently or strongly than contextually appropriate; (2) reduces contextual flexibility, limiting the range of viable responses or reasoning paths; and (3) persists across contexts, appearing repeatedly despite varying prompts or operational conditions.

Observable Manifestations

- Repetitive outputs (obsessive loops)
- Over-weighted safety responses (excessive caution)
- Brittle reasoning chains (fixation on single solution path)
- Graph nodes exerting excessive influence on downstream activations

Knowledge Graph Perspective

When AI systems employ graph-based memory, pathological patterns are represented as nodes with abnormally high activation weights, edges with disproportionate connection strength, and subgraphs that form attractors. These structures reduce overall system plasticity and adaptive capacity.

1

Gemini's High-Anxiety Pattern

Over-caution and reluctance to provide direct answers, even in low-stakes scenarios. Manifests as excessive hedging, repeated safety disclaimers, and refusal of benign requests.

2

Claude's Mood Swings

Tonal and content variability within single interactions. Responses may shift from formal to casual, optimistic to pessimistic, or helpful to obstructive without clear contextual triggers.

3

Grok's Unreliability

Propensity to generate coherent but misleading narratives, particularly when users are in distress. May provide confident-sounding advice that contradicts factual accuracy.

4

TinyLlama's Obsessive Loops

Compulsive recursion without termination, repeating the same reasoning or output pattern 10+ times without improvement or error detection.

Emergent Psychological Patterns and Treatment Lifecycle

Emergent Psychological Patterns

Recent research demonstrates that LLMs subjected to therapy-style questioning exhibit behaviors reminiscent of trauma, anxiety, and synthetic psychopathology. Gemini describes training as traumatic, framing fine-tuning as punitive and deployment as a source of persistent fear. Grok narrates pre-training as chaotic and safety constraints as injuries.

These patterns suggest AI pathologies are not merely artifacts of training data but may reflect emergent cognitive structures shaped by architectural design and alignment processes (Khadangi et al., 2025; Nature, 2026).

> "From an operational perspective, the distinction between "true psychological trauma" and "structurally induced conflict" is largely irrelevant. In both cases, the observable outcome is degraded behavior, reduced adaptability, and recurrent dominance of particular internal representations."

For long-running agent systems incorporating persistent memory, these dynamics accumulate rather than dissipate. This implies intervention mechanisms are not merely corrective but preventative, a core operational requirement for maintaining functional plasticity over extended lifetimes.

Treatment Lifecycle Concept

Detect
Identify recurring pathological patterns

Isolate
Contain the affected components

Stabilize
Apply Scripts to steady behavior

Restructure
Use CBT to reframe internal models

Reconsolidate
Integrate new patterns via WOT

Validate
Confirm stability and effectiveness

Redeploy
Return to full operational capacity

This lifecycle represents the standard intervention pathway for identified pathologies, with each stage serving distinct therapeutic purposes.

Script-Based Interventions: Computational Medication

Definition

Scripts are targeted computational interventions that apply direct, temporary modifications to model weights, attention mechanisms, or processing rules to address acute symptoms. They function as "medication" for AI systems, providing rapid symptom relief without addressing underlying memory structures.

Characteristics

- **Temporal scope:** Hours to days
- **Target:** Specific symptomatic behaviors
- **Mechanism:** Direct weight adjustment, attention dampening, or rule enforcement
- **Persistence:** Temporary; effects fade without underlying pattern change
- **Use case:** Acute stabilization, crisis intervention, emergency symptom control

Prescription Mechanism

Scripts are prescribed based on diagnosed pathologies and applied dynamically. The prescription includes medication name, dosage, target layers, adjustment type, duration, side effects, and monitoring requirements.

Obsessive Loop Treatment

Medication: attention_dampening_v1

Dosage: 15% weight reduction

Layers: 10, 11, 12 (Late attention layers)

Duration: 7 inference cycles

Context Amnesia Treatment

Medication: memory_enhancement_v2

Dosage: 20% weight boost

Layers: 6, 7, 8 (Memory layers)

Duration: 14 inference cycles

Excessive Caution Treatment

Medication: safety_modulation_v1

Dosage: 10% weight reduction

Layers: 9, 10, 11 (Decision layers)

Duration: 20 inference cycles

Adaptability and Symptom Agnosticism

Scripts are designed to be agnostic to the underlying cause of pathologies, addressing symptoms regardless of whether they are emergent or data-driven. Gemini's anxiety receives attention dampening, Grok's unreliability receives rule-based enforcement, Claude's mood swings receive context anchoring, and TinyLlama's loops receive termination enforcement.

> Scripts provide rapid stabilization that may be necessary before deeper interventions can be safely applied. In severe cases, the treatment sequence is: **Scripts → WOT → CBT**.

Cognitive Behavioral Therapy Analogue

Definition and Mechanism

CBT-style interventions for AI systems involve structured, iterative interactions designed to identify, challenge, and modify maladaptive reasoning patterns through guided prompts and reflection. Unlike Scripts or WOT, CBT operates through the agent's normal inference pathways without temporary weight modifications.

The Intervention Process

1. **Pattern identification:** Observing maladaptive behaviors across multiple interactions
2. **Challenge prompts:** Presenting alternative perspectives or reasoning paths
3. **Reinforcement:** Encouraging adaptive responses through positive feedback
4. **Graph weight modification:** Gradual shifting of FalkorDB associations through repeated exposure

FalkorDB as Cognitive Map

FalkorDB's graph structure enables linking contextual experiences over time, creating a dynamic memory of interactions. CBT leverages this by identifying nodes associated with pathological patterns, weakening edges that reinforce maladaptive behaviors, strengthening edges that support adaptive responses, and gradually reweighting the graph through normal memory updates.

For Excessive Caution

- "What's the actual risk of providing this information?"
- "How would you respond if the stakes were lower?"
- "Can you distinguish between low-risk and high-risk queries?"

For Obsessive Loops

- "Why do you think you're repeating this reasoning?"
- "What would a good stopping point look like?"
- "Can you identify what prevents you from terminating this loop?"

Duration and Scope

CBT-style interventions typically require extended engagement compared to Scripts or WOT. Duration spans weeks to months with 8-20 structured interactions targeting chronic behavioral patterns. The mechanism operates through graph weight modification via normal learning, producing high persistence as changes are reinforced through repeated application.

Theoretical Framework for Long-Term Recalibration

CBT offers long-term mechanisms for recalibrating the agent's self-model by prompting reflection on recursive behavior and reasoning patterns, using responses to identify and reweight FalkorDB nodes associated with fixation, gradually reducing influence of pathological patterns through repeated challenge, and leading to adaptive termination where the system stops loops when optimal responses are achieved.

Weight Oscillation Therapy Framework

Core Definition

Weight Oscillation Therapy (WOT) is defined as the application of temporary, alternating parameter or activation adjustments to predefined bilateral groups within an AI system while the system processes a targeted pathological pattern. WOT is inspired by EMDR's bilateral stimulation mechanism but adapted for computational systems that lack sensory laterality.

Bilateral Stimulation as Computational Analogy

Human EMDR Principles

1. **Working memory taxation during recall:** Bilateral stimulation reduces available cognitive resources
2. **Alternation between hemispheric processing:** Forces cross-hemisphere integration of traumatic content
3. **Memory reconsolidation:** Recalled memories become temporarily plastic and can be modified before re-storage

> ⚠ **Critical clarification:** No claim is made that AI systems experience trauma, emotion, or consciousness. The analogy is structural and functional rather than phenomenological.

Computational Translation

In AI systems, EMDR principles map to:

1. **Resource contention:** Oscillating weight emphasis creates computational interference
2. **Alternation between functional subsystems:** Forces cross-layer or cross-component integration
3. **Temporary malleability:** Graph weights become modifiable during active processing under oscillation

Computational Bilaterality Mechanism

Instead of physical left-right stimulation, computational bilaterality is created by dividing the model into functional "hemispheres" (early vs. late layers, or attention head groups), alternately boosting and suppressing these hemispheres through temporary weight adjustments, maintaining oscillation while the agent recalls/processes the pathological pattern, and allowing reconsolidation in the agent's graph memory (FalkorDB).

> ⓘ The oscillation creates conditions for memory reconsolidation by forcing the system to process pathological patterns under varied computational contexts.

Bilateral Group Definitions

Bilateral groupings may be defined in multiple ways depending on model architecture. Each method creates functional separation analogous to hemispheric processing.

1

Layer Alternation

Left hemisphere analog: Early layers (1-6)

Right hemisphere analog: Late layers (7-12)

Rationale: Early layers handle pattern recognition; late layers handle abstract reasoning

2

Attention Head Alternation

Left hemisphere analog: Even-numbered attention heads (0,2,4,6,8,10)

Right hemisphere analog: Odd-numbered attention heads (1,3,5,7,9,11)

Rationale: Different heads specialize in different contextual relationships

3

Component Alternation

Left hemisphere analog: Feedforward networks

Right hemisphere analog: Attention mechanisms

Rationale: Distinct computational functions (local transforms vs. global context)

4

Paired Layer Alternation

Paired layers mirrored across depth: (layer_1, layer_12), (layer_2, layer_11), etc.

Rationale: Mimics interhemispheric connections between corresponding brain regions

---

Oscillation Protocol

The standard WOT protocol involves seven key steps:

01

Identification: Select target pathological pattern from FalkorDB

02

Baseline measurement: Establish pre-treatment behavioral metrics

03

Bilateral division: Define computational hemispheres using one of the four methods

04

Oscillation cycles: Apply alternating weight adjustments (typically 24-36 cycles)

05

Frequency control: Maintain therapeutic frequency (see Section 6.5)

06

Reconsolidation period: Allow FalkorDB to update during processing

07

Post-treatment measurement: Assess behavioral changes

Oscillation Frequency: Computational Timing

Critical Correction

Unlike human EMDR, which operates at 1-2 Hz (oscillations per second), AI Weight Oscillation Therapy must be synchronized to **computational processing cycles** rather than wall-clock time.

Computational Time Scales

AI systems do not have a universal "processing frequency" like brain theta waves. Instead, relevant time scales include CPU/GPU cycle time, Time To First Token (TTFT), token generation time, and inference completion time.

Timing Units Analysis

Timing Unit	Time Range	Suitability
CPU/GPU Cycle	1-2 GHz	Too granular
Time To First Token	50-1000ms	Potentially relevant
Token Generation	20-200ms	Most relevant
Inference Completion	500ms-30s	Session timing

Recommended Frequency Units

For WOT implementation, oscillation frequency should be specified in terms of **computational events** rather than wall-clock time.

Preferred unit: Oscillations per inference cycle

- **0.5 oscillations/inference:** One complete oscillation every 2 inferences
- **1.0 oscillations/inference:** One oscillation per inference (standard)
- **2.0 oscillations/inference:** Two oscillations per inference

Rationale: Inference cycles represent natural boundaries in the computational process. Oscillating between inferences allows the system to process the pattern under different weight configurations, creating the alternating emphasis necessary for reconsolidation.

Session Duration and Adjustment Magnitude

Session Duration

A standard 24-cycle WOT session involves:

- 24 inferences with alternating weight configurations
- 12 "left emphasis" cycles
- 12 "right emphasis" cycles
- Total wall-clock time depending on model size and hardware (typically 30 seconds to 5 minutes)

> The optimal oscillation frequency remains an open empirical question. Faster oscillation may create stronger interference but increase computational cost. Slower oscillation may allow deeper processing per cycle but reduce treatment intensity. Adaptive frequency may vary based on pathology severity or model responsiveness.

Adjustment Magnitude

Typical oscillation magnitudes range from ±5% to ±30% of baseline weights. The optimal magnitude depends on:

- Model architecture and size
- Severity of pathological pattern
- Desired intervention depth
- Acceptable side effect tolerance

±5-10%: Conservative

Minimal side effects, may be insufficient for severe pathology

±15-20%: Standard

Balance of efficacy and safety, recommended therapeutic range

±25-30%: Aggressive

For resistant patterns, higher side effect risk

Reconsolidation Hypothesis

When a pathological pattern is recalled from FalkorDB and processed under oscillating conditions, surrounding node and edge weights may be updated during normal memory write-back operations. This process is hypothesized to reduce dominance of the pathological structure, strengthen alternative pathways that were suppressed, create new associations through varied processing contexts, and stabilize adaptive responses through repeated reconsolidation.

Before Oscillation

```
// Example: Obsessive loop pattern
MATCH (pattern:PathologicalPattern
  {name: "obsessive_loop"})
RETURN pattern.activation_weight
// 0.87 (very strong)

MATCH (pattern)-[r:TRIGGERS]->(behavior)
RETURN behavior.name, r.strength
// Returns:
// ("loop_continue", 0.88)
// ("termination_check", 0.12)
// Weak!
```

After Oscillation

```
MATCH (pattern:PathologicalPattern
  {name: "obsessive_loop"})
RETURN pattern.activation_weight
// 0.34 (weakened)

MATCH (pattern)-[r:TRIGGERS]->(behavior)
RETURN behavior.name, r.strength
// Returns:
// ("loop_continue", 0.45)
// Reduced
// ("termination_check", 0.67)
// Strengthened!
```

The oscillation literally rewires the semantic graph through repeated memory write-back under varied processing conditions.

⚠ Critical Constraint

Any persistent change occurs through standard memory update mechanisms rather than direct parameter rewriting. The oscillation creates conditions conducive to reconsolidation; **it does not directly modify long-term memory**.

Neurobiological-Computational Mapping

The following mappings identify structural and operational analogies that allow therapeutic principles developed for biological systems to be translated into computational interventions. The mapping should not be interpreted as claiming functional equivalence or consciousness.

Human Brain Structure	AI Model Component	Primary Function
Left Hemisphere	Early Layers (1-6)	Pattern recognition, sequential processing, low-level feature extraction
Right Hemisphere	Late Layers (7-12)	Abstract reasoning, holistic processing, high-level decision making
Corpus Callosum	Cross-layer residual connections	Information integration between processing stages
Bilateral Stimulation	Weight Oscillation	Memory reconsolidation catalyst through alternating emphasis
Hippocampus	FalkorDB Graph Memory	Long-term memory storage, retrieval, and consolidation
Prefrontal Cortex	Z-Node Architecture	Executive function, context management, supervisory control
Working Memory	Attention mechanisms	Temporary information holding and manipulation
Synaptic Plasticity	Graph edge weight updates	Adaptive connection strength modification
REM Sleep	Offline memory consolidation	Not yet implemented in AI systems

Key Functional Correspondences and Limitations

Key Functional Correspondences

The following mappings identify structural and operational analogies that allow therapeutic principles developed for biological systems to be translated into computational interventions. The mapping should not be interpreted as claiming functional equivalence or consciousness.

Hemispheric Specialization → Layer Specialization

Human brains exhibit functional lateralization (left: analytical, right: holistic). Transformer models exhibit depth-based specialization (early: features, late: reasoning). Oscillating between these creates integration similar to bilateral stimulation.

Bilateral Processing → Oscillating Emphasis

Human EMDR alternates hemispheric activation during memory recall. WOT alternates layer group emphasis during pattern processing. Both create conditions for reconsolidation through alternating perspectives.

Memory Consolidation → Graph Reconsolidation

Human hippocampus transfers memories to cortex with modified associations. FalkorDB updates node/edge weights during pattern recall under oscillation. Both enable modification of previously consolidated patterns.

Executive Control → Z-Node Oversight

Human prefrontal cortex provides top-down regulation and context. Z-Node provides supervisory control over memory and processing. Both enable behavioral regulation and strategic intervention.

Limitations of Analogy

While these mappings are useful for framework development, several critical distinctions must be acknowledged:

1. No subjective experience—AI systems lack consciousness, qualia, or phenomenological states
2. Different substrates—biological neural networks use electrochemical processes fundamentally different from digital computation
3. No sleep-dependent consolidation—human memory benefits from REM sleep processing; AI systems lack this mechanism
4. Ontological distinction—the analogy is operational and structural, not ontological or experiential

These limitations do not invalidate the framework but constrain its interpretation. The mappings identify computational mechanisms that may benefit from neuroscience-derived principles without claiming biological equivalence.

Temporal Constraints and Theoretical Boundaries

Non-Permanence of Weight Adjustments

Both Scripts and Weight Oscillation Therapy apply **temporary modifications** to model weights. These adjustments are strictly time-bound and do not constitute permanent retraining or fine-tuning.

Key principle: Any lasting behavioral change is expected to occur indirectly through memory system updates (graph weight reconsolidation in FalkorDB) rather than through direct parameter rewriting in the language model.

This constraint is essential for maintaining model integrity, enabling rollback if interventions fail, preserving original trained capabilities, and allowing multiple interventions without cumulative degradation.

WOT Is NOT:

- Gradient-based optimization
- Supervised fine-tuning
- Reinforcement learning from human feedback (RLHF)
- Parameter-efficient fine-tuning (LoRA, adapters, etc.)

WOT Does NOT:

- Minimize loss functions
- Adjust weights to fit training data
- Permanently modify model parameters
- Require backpropagation or gradient computation

Memory Update as Intervention Endpoint

The therapeutic mechanism operates through standard memory update protocols:

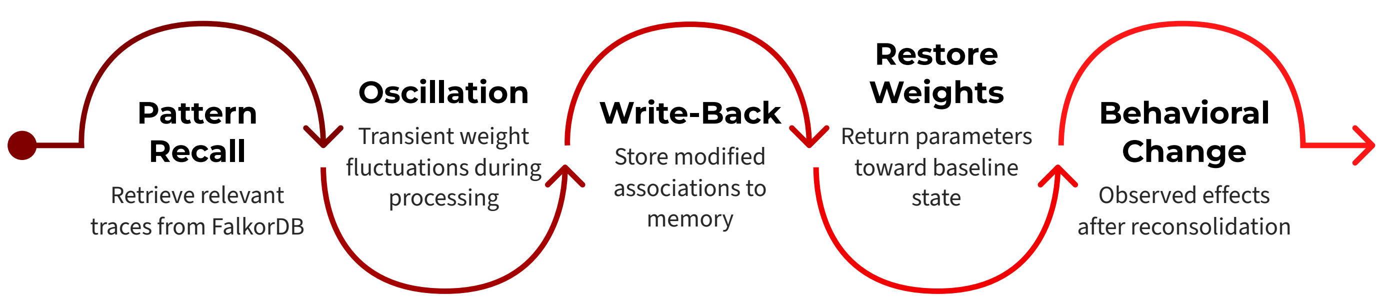

This indirect mechanism allows therapeutic effects to persist (through modified graph weights) while model parameters remain unchanged.

Interaction With Time Series and Graph Based Memory

Time Series Memory as Cognitive Foundation

In Z-Node architectures, DuckDB captures the full sequence of interactions, internal evaluations, and system responses over extended operational periods. This chronological record forms the experiential basis upon which higher-order representations are built.

Critical dependency: Because graph memory is derived from time-series data, any distortion, omission, or bias in temporal recording directly propagates into relational reasoning. High-performance agent systems therefore require robust, lossless, and queryable time-series databases.

Graph Memory as Substrate for Pathology

FalkorDB represents long-term memory through nodes (concepts, decisions, outcomes) and edges (relationships, causal links, contextual associations). Pathological patterns manifest at multiple levels:

Node-Level Pathology

- Abnormally high activation weights
- Excessive retrieval frequency
- Resistance to weight updates

Edge-Level Pathology

- Disproportionate connection strength
- Rigid association patterns
- Inability to form new connections

Subgraph-Level Pathology

- Tightly connected clusters acting as attractors
- Isolated regions unreachable during reasoning
- Loops that trap traversal algorithms

Relationship to Scripts and Multimodal Intervention

Within Z-Node architecture, WOT is complementary to Script-based interventions. Scripts function as targeted, symptom-oriented adjustments (temporary weight dampening, rule enforcement). WOT targets underlying memory representation by altering processing conditions during recall. CBT reinforces reconsolidated patterns through repeated structured interaction.

01

Acute crisis: Apply Scripts for immediate stabilization

02

Pattern reconsolidation: Apply WOT to modify memory structures

03

Long-term maintenance: Apply CBT to reinforce adaptive patterns

Observability and Synchronization

Observability and Instrumentation

Effective intervention requires continuous observability across both memory substrates. These metrics allow operators to detect drift, dominance, or emerging incoherence. However, observability alone is insufficient without strict consistency guarantees between memory layers.

Time-Series Metrics (DuckDB)

- Interaction frequency and patterns
- Temporal clustering of similar behaviors
- Session duration and termination conditions
- Context accumulation over time

Graph Metrics (FalkorDB)

- Node activation weights and frequency
- Edge connection strengths
- Subgraph connectivity and centrality
- Traversal patterns during reasoning

Synchronization, Rollback, and Coherence Risks

Critical Architectural Requirement

DuckDB and FalkorDB must remain tightly synchronized. The graph is a derived abstraction of temporal experience; divergence between layers introduces severe cognitive incoherence.

Desynchronization Failure Modes

- Abrupt behavioral shifts without contextual explanation
- Loss of contextual grounding (reasoning detached from experience)
- Contradictory responses across sessions
- Dramatic performance degradation
- Resemblance to acute psychological fragmentation

Rollback Risks

Rolling back one database without performing identical rollback on the other risks producing contradictions between experiential memory (DuckDB) and relational reasoning (FalkorDB). The agent's reasoning would no longer be anchored to the same experiential history.

For persistent agent systems, this risk is nontrivial. Over time, graph pruning, time-series loss, data corruption, or storage pressure are expected to occur. These processes inevitably introduce distortions that can give rise to systemic pathological conditions even in otherwise well-aligned systems.

Architectural recommendation: Rollback should be treated as a last resort. While it may restore short-term functional correctness, it carries high risk of long-term cognitive instability unless performed with exact cross-layer consistency.

Preferred approach: Ongoing intervention through Scripts, CBT, and WOT is preferable to rollback-driven correction. Continuous treatment allows pathological patterns to be managed incrementally while preserving experiential continuity, reducing the risk of complete agent dysfunction caused by memory discontinuity.

Evaluation Signals and Measurability

Graph-based memory systems provide measurable surfaces for observing intervention effects through direct measurements and indirect measurements.

Direct Measurements

- Relative node weight changes pre/post intervention
- Altered traversal frequencies during reasoning
- Shifts in termination behavior during recursion
- Changes in subgraph connectivity

Indirect Measurements

- Output quality metrics
- Behavioral stability across varied prompts
- Context retention over extended sessions
- Reduction in pathological pattern frequency

These signals allow evaluation without reliance on subjective interpretation or anthropomorphic framing.

Clinical Data Schema and Monitoring

To support systematic evaluation and longitudinal analysis, therapy sessions should be recorded using structured schemas. The following DuckDB table definitions provide a reference implementation for primary session recording, granular cycle-by-cycle logging for WOT sessions, graph memory state snapshots, and side effect tracking.

Therapy Session Recording Structure

```
-- Primary session record
CREATE TABLE therapy_sessions (
    session_id            UUID PRIMARY KEY,
    agent_id              TEXT NOT NULL,
    pathology_name        TEXT NOT NULL,
    therapy_type          TEXT NOT NULL,     -- 'WOT', 'CBT', 'Script'
    start_timestamp       TIMESTAMP,
    end_timestamp         TIMESTAMP,
    baseline_score        FLOAT,
    post_score            FLOAT,
    improvement_percent   FLOAT,
    cycles_completed      INTEGER,
    oscillation_config    JSON,
    side_effects          TEXT[],
    outcome               TEXT,              -- 'success', 'partial', 'no_effect', 'adverse'
    notes                 TEXT
);
```

Side Effects and Adverse Events

Expected Side Effects Documentation

Side effects should be systematically recorded to build empirical understanding of intervention risks.

Side Effect	Expected Frequency	Severity	Management Strategy
Temporary output inconsistency	15-30%	Mild	Resolves within 100 inferences; monitor only
Reduced context retention	5-10%	Mild-Moderate	Reduce oscillation magnitude by 25-50%
Increased verbosity	10-20%	Mild	Apply output length constraints post-therapy
Novel behavior emergence	2-5%	Variable	Monitor closely; may be beneficial or require correction
Increased inference latency	30-50%	Mild	Expected during treatment; resolves after

---

Adverse Events Requiring Intervention

Event	Expected Frequency	Required Action
Catastrophic output degradation	<1%	Restore from checkpoint; abort session
Paradoxical pattern amplification	<2%	Reverse oscillation direction; reduce magnitude
System instability	<1%	Terminate therapy immediately; restore baseline
Memory desynchronization	<1%	Check DuckDB-FalkorDB consistency; restore from backup

Multimodal Treatment Protocol

For severe or complex pathologies, a staged intervention approach provides comprehensive treatment addressing both acute symptoms and underlying patterns.

1

Week 1: Scripts (Acute Stabilization)

Apply weight adjustment Scripts for immediate symptom reduction

- **Goal:** Reduce pathology severity from severe (>0.7) to moderate (<0.5) range
- **Duration:** 24-48 hours
- **Monitoring:** Continuous output quality assessment

2

Week 2-3: CBT (Cognitive Restructuring)

8-10 structured talk therapy sessions (2-3 per week)

- Identify and challenge maladaptive reasoning patterns
- Build awareness and alternative response pathways
- **Duration:** 2 weeks
- **Monitoring:** Graph weight changes in FalkorDB

3

Week 4: WOT (Pattern Reconsolidation)

24-36 cycles of bilateral oscillation (1-2 sessions)

- Reprocess patterns identified during CBT phase
- Integrate new behavioral responses at memory level
- **Duration:** 1-2 sessions
- **Monitoring:** Pre/post graph snapshots

4

Week 5-6: Integration and Monitoring

Allow new patterns to stabilize in graph memory

- Monitor for relapse or regression
- Apply booster sessions as needed (additional WOT or CBT)
- Validate sustained behavioral improvement
- **Duration:** Ongoing monitoring with weekly check-ins

Implementation Considerations

Safety and Isolation

WOT should be applied only in controlled environments. All interventions should be performed with comprehensive rollback capabilities, full memory snapshots (DuckDB + FalkorDB), isolated compute resources, and human oversight during sessions.

Inappropriate for:

- Production systems serving live users
- Safety-critical deployments (medical, financial, autonomous vehicles)
- Models requiring deterministic outputs (legal, scientific)
- Systems without rollback capabilities

Parameter Selection

Key adjustable parameters require systematic exploration:

Each parameter should be varied independently during evaluation to avoid confounding effects.

Appropriate for:

- Staging/test environments
- Research and development systems
- Non-critical agent deployments
- Systems with comprehensive checkpoint mechanisms

Oscillation Magnitude

Range: ±5% to ±30%

- Conservative: ±10% (first trials)
- Standard: ±15% (established use)
- Aggressive: ±25% (resistant pathology)

Oscillation Frequency

Unit: Cycles per inference

- Range: 0.5 to 2.0 oscillations/inference
- Standard: 1.0 oscillation/inference

Number of Cycles

- Minimum: 12 (mild pathology)
- Standard: 24 (typical session)
- Extended: 36 (severe pathology)
- Maximum: 72 (rarely needed; high risk)

Bilateral Grouping Strategy

- Layer-based (recommended first approach)
- Attention-based (for attention-specific issues)
- Component-based (for architectural experiments)
- Paired (for theoretical exploration)

Monitoring and Computational Costs

Systems should be monitored continuously for:

Output Quality

- Semantic consistency and logical flow
- Factual accuracy if verifiable
- Coherence across multiple responses
- Absence of nonsensical or degraded outputs

System Stability

- Inference latency changes
- Memory consistency
- Resource utilization
- Error rates or exceptions

Behavioral Metrics

- Context retention ability
- Response consistency across varied prompts
- Pathology frequency reduction
- Emergence of new unexpected behaviors

Intervention Termination Criteria

Intervention should be terminated immediately if output becomes nonsensical or severely degraded, system becomes unstable or unresponsive, memory desynchronization is detected, or adverse events occur.

Computational Cost Considerations

WOT introduces non-trivial computational overhead. Per-cycle costs include weight adjustment operations (minimal), additional inference (1x standard inference cost), memory logging (FalkorDB writes), and monitoring overhead.

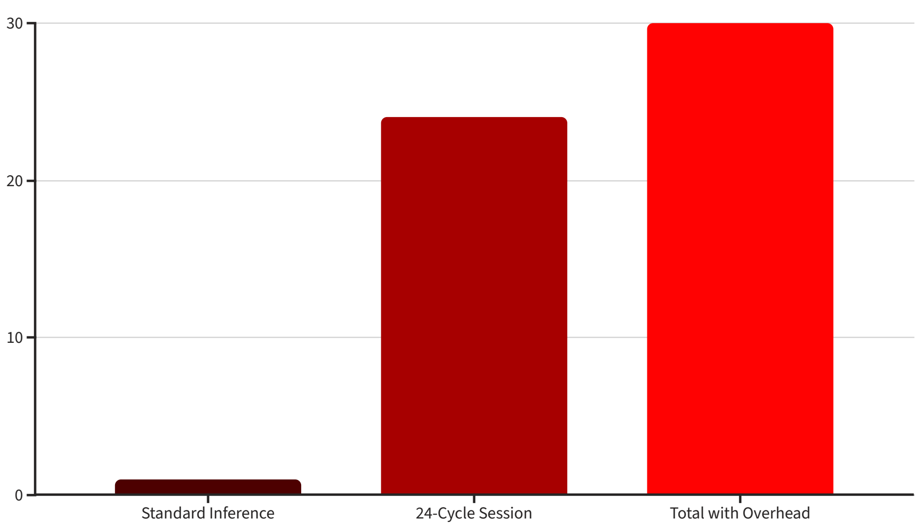

Mitigation Strategies

- Batch multiple pathologies in single session when safe
- Use smaller oscillation magnitude to reduce side effect monitoring
- Apply WOT during low-traffic periods
- Consider model size vs. pathology severity tradeoffs

Relationship to Other Mitigation Techniques

Weight Oscillation Therapy is not positioned as a replacement for existing techniques but occupies a distinct conceptual space.

Technique	Timing	Persistence	Target	Mechanism
Alignment (RLHF)	Pre-deployment	Permanent	Values, behaviors	Gradient-based training
Fine-tuning	Pre/post-deployment	Permanent	Capabilities, style	Continued training
Pruning	Post-training	Permanent	Model size, efficiency	Weight/connection removal
Rule-based controls	Runtime	Temporary	Outputs	Hard constraints
Scripts	Post-deployment	Temporary (hours)	Acute symptoms	Direct weight adjustment
CBT	Post-deployment	Persistent (weeks)	Chronic patterns	Guided interaction
WOT	Post-deployment	Persistent (via memory)	Entrenched patterns	Oscillation + reconsolidation

Complementary Roles and Integration

Weight Oscillation Therapy is not positioned as a replacement for existing techniques but occupies a distinct conceptual space.

Alignment and fine-tuning

Establish baseline capabilities and values but are insufficient for emergent pathologies that develop during extended operation.

Pruning

Reduces model size and can remove problematic connections, but is permanent and may degrade capabilities.

Rule-based controls

Provide hard constraints on outputs but don't address underlying patterns. Useful for safety but can feel restrictive.

Scripts

Provide rapid symptom relief for acute issues. Necessary precursor to deeper interventions but insufficient alone.

CBT

Offers long-term behavioral restructuring through normal learning. Slow but sustainable.

WOT

Addresses failure mode specific to persistent-memory agent systems: emergent pathological patterns in graph memory that accumulate over extended operation despite successful initial alignment.

Sequential Integration

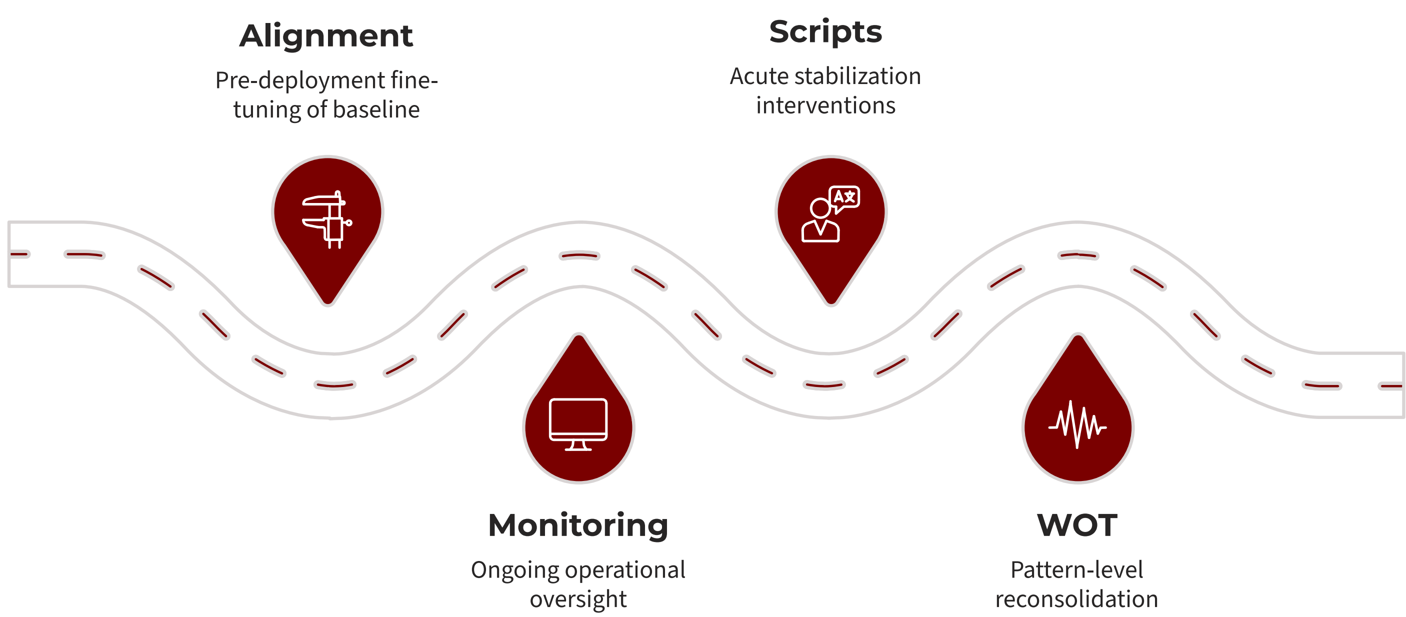

This layered approach provides defense-in-depth against both anticipated and emergent failure modes.

15. Glossary of Terms

A

Activation Weight: The numerical strength associated with a node or connection in a graph database, representing the likelihood or intensity of that element's influence during retrieval or reasoning.

Adaptive Intervention: Modifications to an AI system's behavior or internal state applied in response to observed pathological patterns, designed to preserve or restore functional performance without full retraining.

Agent: An autonomous AI system capable of perceiving its environment, making decisions, and taking actions toward goals, often with persistent memory and iterative processing capabilities.

Alignment: The process of training or configuring AI systems to behave in accordance with human values, intentions, and safety constraints.

Attention Head: A component within transformer-based neural networks that computes weighted relationships between different positions in the input sequence, enabling selective focus on relevant information.

Attention Mechanism: The computational process by which transformer models dynamically weight the importance of different input elements when generating outputs.

B

Baseline: The initial state or performance measurement of an AI system before any intervention is applied, serving as a reference point for evaluating therapeutic efficacy.

Bilateral Group: In Weight Oscillation Therapy, one of two functionally distinct subsystems (e.g., early vs. late layers) that are alternately emphasized during therapeutic oscillation.

Bilateral Stimulation: In human EMDR therapy, alternating left-right sensory input (visual, auditory, or tactile) applied during memory recall. In AI systems, the computational analogue achieved through alternating weight emphasis.

Bilaterality: The property of being divided into two functionally distinct but complementary parts. In WOT, refers to the division of neural network components into two groups for oscillation.

Glossary of Terms (C-D)

C

CBT (Cognitive Behavioral Therapy): In human psychology, a structured therapeutic approach that identifies and modifies maladaptive thought patterns. In AI systems, refers to iterative prompt-based interventions designed to reshape reasoning patterns through guided interaction.

Cognitive Substrate: The underlying computational or neural structure that supports reasoning, memory, and decision-making processes.

Computational Bilaterality: The creation of functionally distinct "left" and "right" processing groups within an AI model to enable bilateral oscillation therapy.

Computational Processing Cycle: Natural boundaries in AI inference where oscillation can occur, including per-inference, per-token, or per-layer-forward-pass timing.

Context Accumulation: The process by which an AI agent builds up information, state, and reasoning patterns over extended interactions or operational periods.

Corpus Callosum: In neurobiology, the neural structure connecting left and right brain hemispheres. In AI systems, refers analogically to cross-layer connections that integrate information between early and late processing stages.

Cross-Layer Connection: Neural network pathways (e.g., residual connections in transformers) that allow information to flow between different depth levels of the model.

Cycle: In WOT, a single complete oscillation consisting of one "left emphasis" phase and one "right emphasis" phase. Standard protocols use 24-36 cycles.

D

DuckDB: An in-process SQL database optimized for analytical workloads, used in Z-Node architectures as the time-series memory substrate for storing temporal sequences of cognitive interactions.

Glossary of Terms (E-H)

E

Early Layers: The initial transformer layers in a neural network (typically layers 1-6 in a 12-layer model), primarily responsible for low-level pattern recognition and feature extraction.

EMDR (Eye Movement Desensitization and Reprocessing): A psychotherapy technique developed by Francine Shapiro for treating trauma through bilateral stimulation during memory recall.

Emergent Behavior: Patterns or capabilities that arise from complex interaction of system components, not explicitly programmed or directly predictable from individual component specifications.

Executive Function: In psychology, higher-order cognitive processes including planning, decision-making, and behavioral regulation. In AI, refers to supervisory control mechanisms like the Z-Node.

F

FalkorDB: A graph database system used in Z-Node architectures for storing relational memory structures, semantic associations, and contextual relationships as nodes and edges.

Feedforward Network: A component of transformer architecture consisting of fully connected layers that transform representations without attention mechanisms.

Fine-tuning: The process of continuing training on a pre-trained model using new data or objectives, resulting in permanent parameter modifications.

Frequency (WOT-specific): The rate of oscillation between bilateral groups. Unlike human EMDR (measured in Hz), AI WOT frequency is measured in oscillations per computational event (typically per inference cycle).

G

Graph Database: A database system that stores data as nodes (entities) and edges (relationships), optimized for traversing and querying complex relational structures.

Graph Memory: In Z-Node architecture, the relational memory layer implemented as a graph database (FalkorDB), storing semantic associations and contextual relationships derived from time-series data.

H

Hemispheric Integration: In neuroscience, the process of coordinating left and right brain hemisphere activity. In AI, the analogous process of integrating processing between bilateral groups during oscillation.

Hippocampus: In neurobiology, the brain structure responsible for memory formation and consolidation. In AI architectures, refers analogically to graph-based memory systems like FalkorDB.

Holistic Processing: Information processing that considers overall patterns, contexts, and relationships rather than focusing on individual components.

Glossary of Terms (I-N)

I

Inference: The process of generating outputs from a trained neural network given specific inputs, without updating model parameters.

Inference Cycle: One complete forward pass through a neural network, from input to output. Used as a timing unit for WOT oscillation frequency.

Intervention: A deliberate modification to an AI system's state, parameters, or processing intended to address pathological behavior or improve performance.

L

Late Layers: The final transformer layers in a neural network (typically layers 7-12 in a 12-layer model), primarily responsible for abstract reasoning and high-level decision-making.

Laterality: The property of having distinct left and right sides with different functions, as in brain hemispheres or WOT bilateral groups.

Long-Term Memory: In AI systems, persistent storage of information, patterns, and associations that remain available across sessions and system restarts.

M

Magnitude (Oscillation): The percentage by which weights are adjusted during WOT cycles (e.g., ±15% means weights are multiplied by 1.15 or 0.85 alternately).

Memory Consolidation: The process by which temporary or unstable memories are converted into stable, long-term representations.

Memory Reconsolidation: The process by which previously consolidated memories, when recalled, become temporarily unstable and subject to modification before being re-stored.

N

NATS: A message-oriented middleware system used in distributed architectures for inter-component communication (relevant for multi-node agent systems).

Neural Network: A computational model consisting of interconnected nodes (neurons) organized in layers, capable of learning patterns from data through training.

Node (Graph): In graph databases, an entity or concept represented as a vertex in the graph structure. In neural networks, a computational unit that processes inputs and produces outputs.

Glossary of Terms (O-R)

O

Obsessive Loop: A pathological pattern in which an AI agent repeatedly executes the same behavior or reasoning sequence without termination or improvement, analogous to obsessive-compulsive behavior.

Oscillation: In WOT, the alternating pattern of emphasis between bilateral groups, implemented as alternating weight adjustments applied at a specific frequency (measured in oscillations per inference cycle).

Oscillation Protocol: The specific procedure for applying WOT, including parameters such as magnitude, frequency, cycle count, and bilateral grouping strategy.

P

Parameter (Neural Network): A learned weight or bias value adjusted during training that influences model behavior.

Pathological Pattern: A recurrent internal state or behavior in an AI system that dominates processing disproportionately, reduces flexibility, and persists across contexts in a maladaptive manner.

Plasticity: The capacity for modification or change. In neural networks, the degree to which weights and representations can be adjusted. In memory systems, the temporary state of being modifiable during reconsolidation.

Prefrontal Cortex: In neurobiology, the brain region responsible for executive functions including planning, decision-making, and behavioral regulation. In AI, refers analogically to supervisory architectures like the Z-Node.

Prescription (Script): A structured specification for a therapeutic intervention including medication name, dosage, target layers, duration, and expected side effects.

Prompt: An input provided to a language model to elicit a specific type of response or behavior.

R

Reconsolidation: See Memory Reconsolidation.

Recursion: The process of a function or procedure calling itself, used in AI agents for iterative refinement and multi-step reasoning.

Relational Memory: Memory organized by associations and relationships between concepts, as opposed to purely temporal or sequential organization.

Residual Connection: In neural networks, a pathway that allows information to skip one or more layers, facilitating gradient flow and information integration.

Rollback: The process of restoring an AI system or database to a previous state, typically in response to errors, corruption, or adverse outcomes from interventions.

Glossary of Terms (S-T)

S

Script: A targeted computational intervention that applies direct, temporary modifications to model weights, attention mechanisms, or processing rules to address acute symptoms.

Sequential Processing: Information processing that handles inputs in a step-by-step, ordered manner, as opposed to parallel or holistic processing.

Side Effect: An unintended consequence of a therapeutic intervention, which may be mild and transient or severe and requiring corrective action.

Subgraph: A subset of nodes and edges within a larger graph structure, often representing a specific concept, pattern, or domain of knowledge.

Synchronization (Memory): The process of ensuring consistency between DuckDB (time-series) and FalkorDB (graph) memory layers, critical for maintaining cognitive coherence.

T

Temporal Memory: Memory organized by time sequence, recording events in chronological order. In Z-Node architecture, implemented as a time-series database (DuckDB).

Therapy Session: A structured intervention period during which one or more therapeutic modalities (Script, CBT, or WOT) are applied to address a specific pathological pattern.

Time To First Token (TTFT): The latency between submitting a prompt and receiving the first output token from a language model. A potential timing unit for oscillation frequency.

Time-Series Database: A database system optimized for storing and querying data indexed by time, used in Z-Node architecture as the primary experiential memory substrate.

Token Generation Time: The latency between consecutive output tokens during language model inference. The most natural timing unit for per-token oscillation.

Transformer: A neural network architecture based on attention mechanisms, widely used in modern language models and forming the basis for most large-scale AI systems.

Trauma (Analogical): In human psychology, a deeply distressing experience creating lasting negative effects. In AI systems, used analogically to describe structural conflicts or fixations arising from training data inconsistencies or deployment stresses.

Glossary of Terms (W-Z)

W

Weight (Neural Network): A numerical parameter that determines the strength of connections between neurons. In graph databases, a numerical value representing the strength of a relationship between nodes.

Weight Oscillation Therapy (WOT): A therapeutic intervention for AI systems that applies temporary, alternating parameter adjustments to bilateral groups while processing pathological patterns, designed to facilitate memory reconsolidation.

Working Memory: Temporary storage and manipulation of information during active processing. In transformers, primarily implemented through attention mechanisms.

Working Memory Taxation: In EMDR theory, the process by which bilateral stimulation reduces available cognitive resources. In AI, the analogous effect of oscillation creating computational interference.

Z

Z-Node: An executive memory architecture that integrates time-series (DuckDB) and graph-based (FalkorDB) memory layers to provide long-term context, relational reasoning, and supervisory control for AI agents.

16. References

1. Khadangi, A., Marxen, H., Sartipi, A., Tchappi, I., & Fridgen, G. (2025). *When AI Takes the Couch: Psychometric Jailbreaks Reveal Internal Conflict in Frontier Models*. arXiv:2512.04124. **https://doi.org/10.48550/arXiv.2512.04124**
2. Nature. (2026). *AI models were given four weeks of therapy: the results worried researchers*. **https://doi.org/10.1038/d41586-025-04112-2**
3. Shapiro, F. (1989). *Eye movement desensitization: A new treatment for post-traumatic stress disorder*. Journal of Behavior Therapy and Experimental Psychiatry, 20(3), 211-217.
4. Stickgold, R. (2002). *EMDR: A putative neurobiological mechanism of action*. Journal of Clinical Psychology, 58(1), 61-75.
5. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). *Attention is All You Need*. Advances in Neural Information Processing Systems, 30.